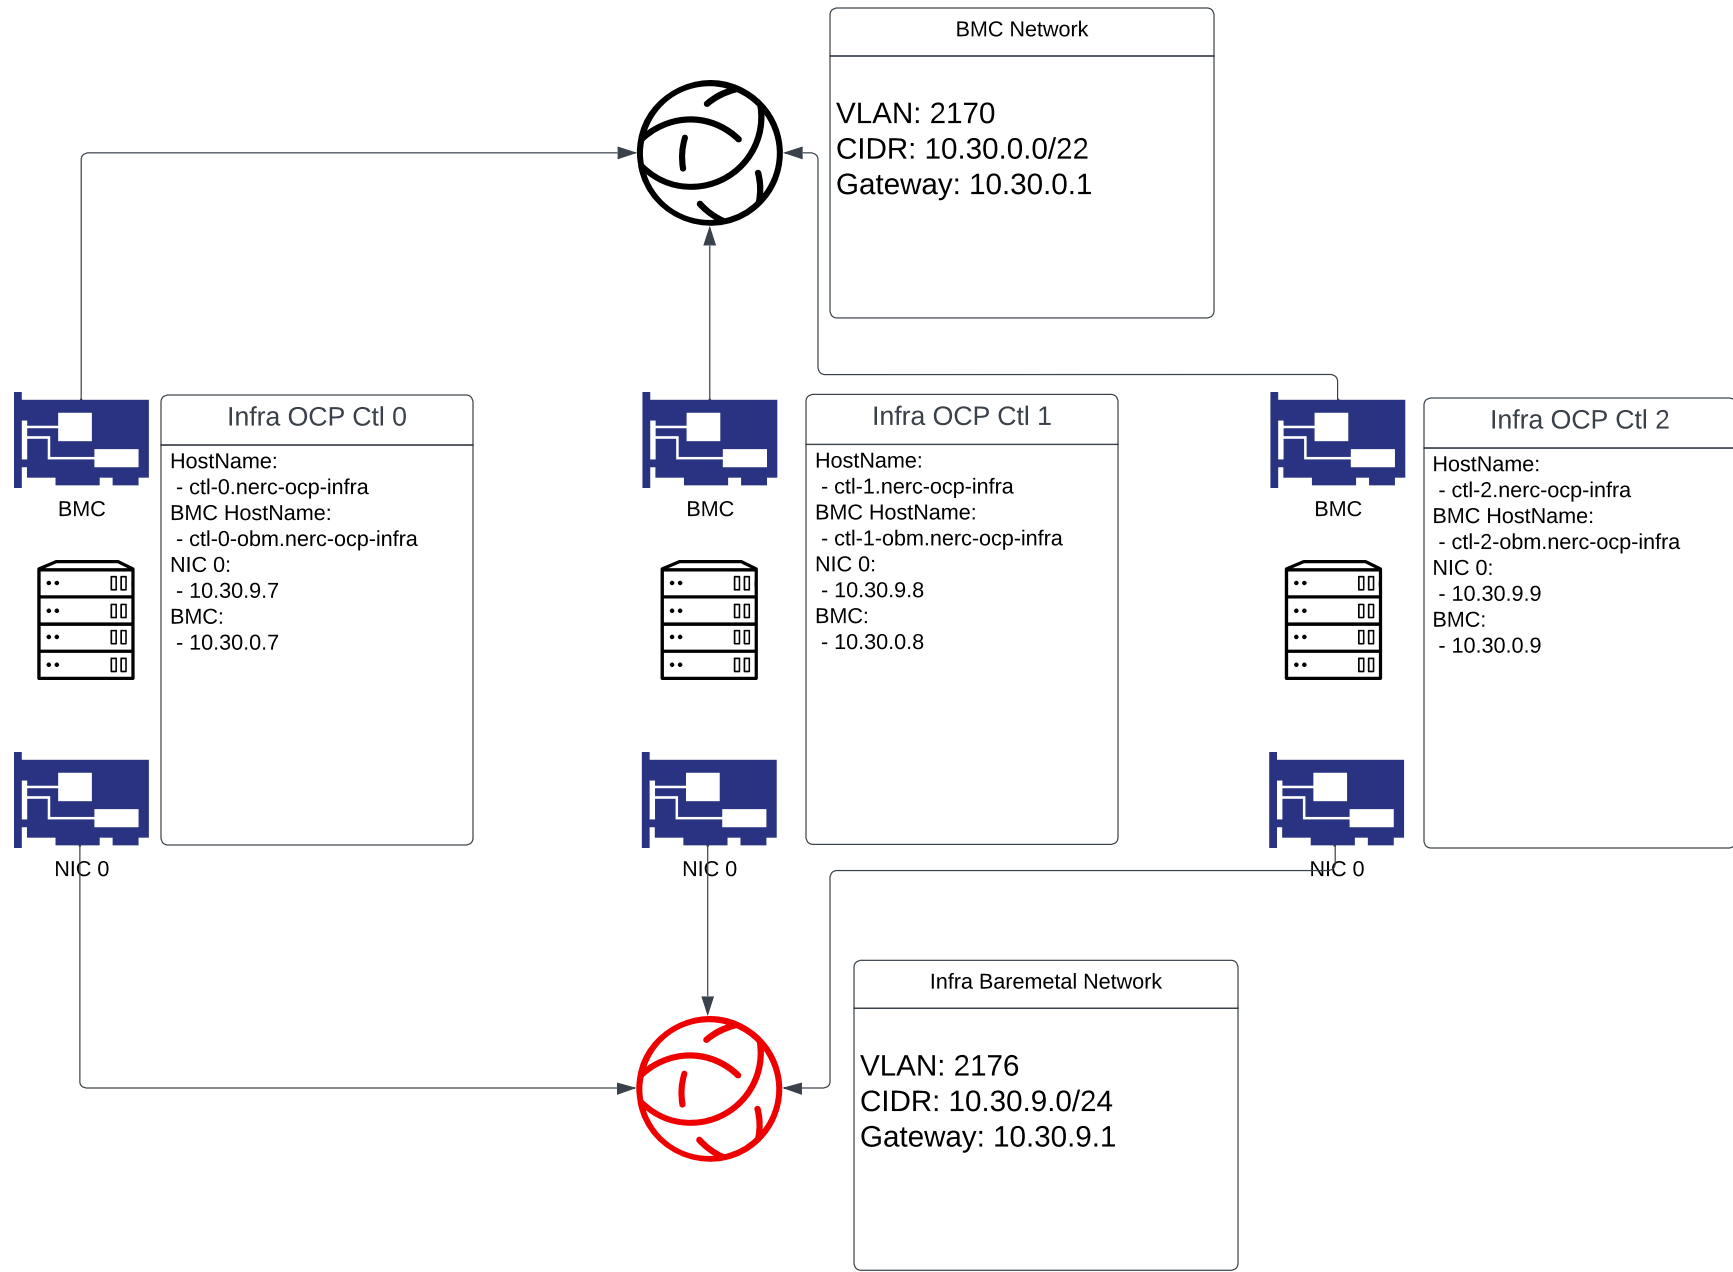








Infra Cluster



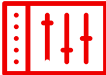
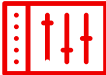

Infra Cluster	
	<div>Details</div> <div>URL: https://console-openshift-console.apps.nerc-ocp-infra.rc.fas.harvard.edu/ API: https://api.nerc-ocp-infra.rc.fas.harvard.edu:6443</div>

Cluster Applications	
ACM	URL: https://multicloud-console.apps.nerc-ocp-infra.rc.fas.harvard.edu/ API: https://api.nerc-ocp-infra.rc.fas.harvard.edu:6443
ArgoCD	URL: https://openshift-gitops-server-openshift-gitops.apps.nerc-ocp-infra.rc.fas.harvard.edu/
Grafana	URL: https://grafana-openshift-monitoring.apps.nerc-ocp-infra.rc.fas.harvard.edu/
Prometheus	URL: https://prometheus-k8s-openshift-monitoring.apps.nerc-ocp-infra.rc.fas.harvard.edu/
Vault	URL: https://vault-ui-vault.apps.nerc-ocp-infra.rc.fas.harvard.edu/

Infra Cluster



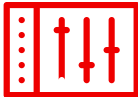
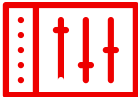
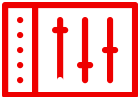
Control Plane








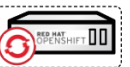

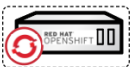



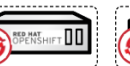

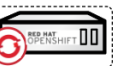






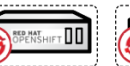


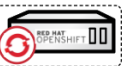














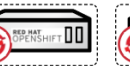
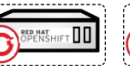
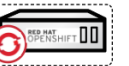
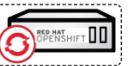


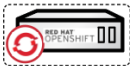



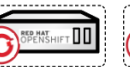
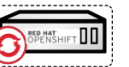


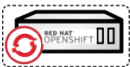



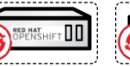








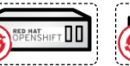








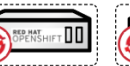
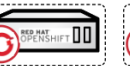
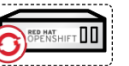
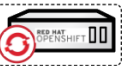


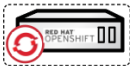
Cluster name: nerc-ocp-infra
Environment: infrastructure/ACM
Control Nodes: 3 (Scheduleable)
Worker Nodes: 0

Production Cluster

Control Plane
















Cluster name: nerc-ocp-prod
Environment: Production
Control Nodes: 3
Worker Nodes: 86




Test Cluster

Cluster name: nerc-ocp-test
Environment: test
Control Nodes: 3
Worker Nodes: 10



Control Plane



Cluster VLAN: 2172
Cluster CIDR: 10.30.6.0/23
Gateway: 10.30.6.1



Storage VLAN: 2173
Storage CIDR: 10.30.10.0/23
Gateway: 10.30.10.1



Cluster Limits

Theoretical limits:

Based on the amount of memory and number of cpu cores available in the control plane

the production cluster should be able to handle over 500 worker nodes.

Cluster Limits

Tested limits:

Red Hat has previously installed and operated a bare metal cluster with 117 worker nodes.

For more info on scaling process and recommendations see:

https://github.com/OCF-on-NERC/docs/blob/main/cluster_scaling_and_load_testing.md